

Compression vs. Perplexity & Throughput
(Qwen2-1.5B — lower size with better metric is ideal)

