

Per-Token Latency by Method × Scope
(AWQ lowest latency; RTN/SmoothQuant higher due to mixed-precision overhead)

